

# COMP3702/COMP7702 Artificial Intelligence (Semester 2, 2020)

## Assignment 4: Reinforcement learning in LASERTANK

### Key information:

- **Due: 5pm, Monday 16 November**
- This assignment will assess your skills in developing reinforcement learning algorithms.
- Assignment 4 contributes 15% to your final grade.
- This assignment consists of two parts: (1) programming and (2) a quiz.
- This is an individual assignment.
- Both code and quiz are to be submitted via Gradescope (<https://www.gradescope.com/>). You can find instructions on how to register for the COMP3702/COMP7702 Gradescope site on Blackboard.
- Your program (Part 1, 60/100) will be graded using the Gradescope code autograder, using the testcases in the support code provided at <https://gitlab.com/3702-2020/assignment-4-support-code>.
- Your quiz answers (Part 2, 40/100) are to be answered directly in Gradescope using the Online Assignment tool (not a document upload).

### The LASERTANK AI Environment

LASERTANK is an open source computer puzzle game requiring logical thinking to solve a variety of levels, originally developed by Jim Kindley in 1995. In LASERTANK, a player controls a tank vehicle which can move forwards, turn clockwise and counterclockwise, and shoot a laser which can be used to interact with special map elements. The game takes place in a 2D map divided into grid cells. The player must navigate from a **starting position** to the **flag** in as few moves as possible while avoiding “game over” conditions. The game is over if the player moves into a dangerous cell (i.e. water cells or any cell in the line of fire area of an anti-tank).

The LASERTANK codebase comes as part of your assignment support code; while a high-level description is provided in a separate LASERTANK AI Environment document, which you can find on Blackboard.

### LASERTANK as a reinforcement learning problem

In this assignment, you will write the components of a program to play LASERTANK, with the objective of finding a high-quality solution to the problem using various reinforcement learning algorithms. This assignment will test your skills in defining a reinforcement learning algorithm for a practical problem and your understanding of key algorithm features and parameters.

### What is provided to you

We will provide supporting code in Python only, in the form of:

1. `laser_tank.py`, This file contains a class representing Laser Tank game map. This is similar to the file provided in Assignment 1, but has a number of important distinctions. Importantly, the instance method `apply_move(self, move)` applies a *non-deterministic action* to the Laser Tank game map, changing it's state. When the `LaserTankMap.MOVE_FORWARD` action is selected, the probability of ending in the square directly ahead is given by `t_success_prob` (an instance variable of the `LaserTankMap` object). All other directions (forward-left, forward-right, left, right and no-change) have equal probability.
2. `policy_visualiser.py`, An animated version of tester which shows each step your agent takes in a simulated episode.

3. A tester script `tester.py`, This will indicate the average reward received by your final learned policy over multiple simulated episodes (50) compared to the benchmark.
4. Testcases to test and evaluate your solution
5. A solver file template

See the support code `README.md` for more information. The support code can be found at: <https://gitlab.com/3702-2020/assignment-4-support-code>. Autograding of code will be done through Gradescope, so that you can test your submission and continue to improve it based on this feedback — you are strongly encouraged to make use of this feedback.

## Your assignment task

Your task is to develop two reinforcement learning algorithms for computing paths (series of actions) for the agent (i.e. the Laser Tank), and to answer a quiz on your algorithms' performance. You will be graded on both your submitted **program (Part 1, 60%)** and the **quiz (Part 2, 40%)**. These percentages will be scaled to the 15% course weighting for this assessment item.

The provided support code provides a generative LASERTANK environment, and your first task is to submit code implementing both of the following reinforcement learning algorithms:

1. Q-learning
2. SARSA

**The support code will include methods allowing the autograder to verify and evaluate each algorithm separately.** The testcase files indicate which algorithm should be applied to them. After you have implemented and tested the algorithms above, you are to complete the questions listed in the section "Part 2 - The Quiz" and submit them directly into Gradescope's Online Assignment tool (nb. *not* a document upload; this differs from Assignments 1 and 2).

More detail of what is required for the programming and quiz parts are given below.

## Part 1 — The programming task

Your program will be graded using the Gradescope autograder, using the COMP3702 testcases in the support code provided at <https://gitlab.com/3702-2020/assignment-4-support-code>.

### Interaction with the testcases and autograder

We now provide you with some details explaining how your code will interact with the testcases and the autograder (with special thanks to Nick Collins for his efforts making this work seamlessly, yet again!).

First, note that the Assignment 4 version of the class `LaserTankMap` (in `laser_tank.py`) differs in that `t_success_prob`, `move_cost`, `collision_cost`, `game_over_cost` and `goal_reward` are now private variables, and are generated randomly upon construction using the initial seed.

Implement your solution using the supplied `solver.py` template file. You should implement your solution by filling in the following method stubs:

- `train_q_learning()`
- `train_sarsa()`
- `get_policy()`

You may add to the `__init__` method if required, and can add additional helper methods and classes (either in `solver.py` or in files you create) if you wish. To ensure your code is handled correctly by the autograder, you should avoid using any try-except blocks in your implementation of the above methods (as this can interfere with our time-out handling). Refer to the documentation in `solver.py` for more details.

### Grading rubric for the programming component (total marks: 60/100)

For marking your programs, we will use six different testcases to verify and evaluate each of the two different algorithm implementations you submit. Marks will be allocated according to the following rules:

- Each algorithm implementation is graded based on performance in each testcase separately.
- *Solving* a testcase means finding a *path* within the given time limit and cost bound (time limits and cost benchmarks are given in each test case file).
- If your code solves a testcase for a given algorithm, it will receive 10 marks. *Approximately solving* a testcase means finding a *sub-optimal solution path*, which is more expensive than the cost benchmark, within the given time limit.
- Approximate solutions are penalised in proportion to how far they are below the cost bound at a rate of 1 mark for each unit of cost it is below the provided benchmark cost, to a minimum of 0 marks.

Students in COMP3702 and COMP7702 face the same grading rubric in this assignment.

## Part 2 — The quiz

The quiz tests your understanding of the methods you have used in your code, and contributes 40/100 of your assignment mark.

**Question 1.** Q-learning is closely related to the value iteration algorithm for Markov decision processes.

- (5 marks) Describe two key similarities between Q-learning and value iteration.
- (5 marks) Give one key difference between Q-learning and value iteration.

For Questions 2, 3 and 4, consider testcase `q-learn_t1.txt/sarsa_t1.txt` (the testcases are the same apart from solving method), and compare Q-learning and SARSA.

**Question 2.**

- (5 marks) With reference to Q-learning and SARSA, explain the difference between off-policy and on-policy reinforcement learning algorithms.
- (5 marks) How does the difference between off-policy and on-policy algorithms affect the way in which Q-learning and SARSA solve testcase `q-learn_t1.txt/sarsa_t1.txt`. Give an example of an expected difference between the way the algorithms learn a policy.

For questions 3 and 4, you are asked to plot the solution quality at each iteration, as given by the 50-step moving average reward received by your learning agent. At time step  $t$ , the 50-step moving average reward is the average reward earned by your learning agent in the iterations  $[t - 50, t]$ , **including episode restarts. If the Q-values imply a poor quality policy, this value will be low. If the Q-values correspond to a high-value policy, the 50-step moving average reward will be higher. We are using a moving average here because the reward is received only occasionally and there are sources of randomness in the transitions and the exploration strategy.**

**Question 3.**

- (5 marks) Plot the quality of the policy learned by Q-learning in testcase `q-learn_t1.txt` against iteration number for three different **fixed** values of the `learning_rate` (which is called  $\alpha$  in the lecture notes and in many texts and online tutorials), as given by the 50-step moving average reward (**i.e. for this question, do not adjust  $\alpha$  over time, rather keep it the same value throughout the learning process**). Your plot should display the solution quality up to an iteration count where the performance stabilises, with a minimum of 2000 iterations (note the policy quality may still be noisy, but the algorithm's performance will stop increasing and its average quality will level out).
- (5 marks) With reference to your plot, comment on the effect of varying the `learning_rate`.

**Question 4.**

- (5 marks) Plot the quality of the learned policy against iteration number under Q-learning and SARSA in testcase `q-learn_t1.txt/sarsa_t1.txt`, as given by the 50-step moving average reward. Your plot should display the solution quality up to an iteration count where the performance of both algorithms stabilise, with a minimum of 2000 iterations.
- (5 marks) With reference to your plot, compare the learning trajectory of the two algorithms, and their final solution quality. Discuss the way the solution quality of Q-learning and SARSA change as they learn to solve the testcase, both as they learn and once they have stabilised.

## Academic Misconduct

The University defines Academic Misconduct as involving “a range of unethical behaviours that are designed to give a student an unfair and unearned advantage over their peers.” UQ takes Academic Misconduct very seriously and any suspected cases will be investigated through the University's standard policy (<https://ppl.app.uq.edu.au/content/3.60.04-student-integrity-and-misconduct>). If you are found guilty, you may be expelled from the University with no award.

It is the responsibility of the student to ensure that you understand what constitutes Academic Misconduct and to ensure that you do not break the rules. If you are unclear about what is required, please ask.

It is also the responsibility of the student to take reasonable precautions to guard against unauthorised access by others to his/her work, however stored in whatever format, both before and after assessment.

In the coding part of this assignment, you are allowed to draw on publicly-accessible resources, but you must make reference or attribution to its source, by doing the following:

- All blocks of code that you take from public sources must be referenced in adjacent comments in your code.
- Please also include a list of references indicating code you have drawn on in your `solution.py` docstring.

**However, you must not show your code to, or share your code with, any other student under any circumstances. You must not post your code to public discussion forums (including Piazza) or save your code in publicly accessible repositories (check your security settings). You must not look at or copy code from any other student.**

All submitted files (code and report) will be subject to electronic plagiarism detection and misconduct proceedings will be instituted against students where plagiarism or collusion is suspected. The electronic plagiarism detection can detect similarities in code structure even if comments, variable names, formatting etc. are modified. If you collude to develop your code or answer your report questions, you will be caught.

For more information, please consult the following University web pages:

- Information regarding Academic Integrity and Misconduct:
  - <https://my.uq.edu.au/information-and-services/manage-my-program/student-integrity-and-conduct/academic-integrity-and-student-conduct>
  - <http://ppl.app.uq.edu.au/content/3.60.04-student-integrity-and-misconduct>
- Information on Student Services:
  - <https://www.uq.edu.au/student-services/>

## Late submission

Students should not leave assignment preparation until the last minute and must plan their workloads to meet advertised or notified deadlines. It is your responsibility to manage your time effectively.

Late submission of the assignment will **not** be accepted. Unless advised, assessment items received after the due date will receive a zero mark unless you have been approved to submit the assessment item after the due date.

In the event of exceptional circumstances, you may submit a request for an extension. You can find guidelines on acceptable reasons for an extension here <https://my.uq.edu.au/information-and-services/manage-my-program/exams-and-assessment/applying-extension> All requests for extension must be submitted on the UQ Application for Extension of Progressive Assessment form at least 48 hours prior to the submission deadline.